

# Laboratorio 7

## Task 1 - Práctica

---

### Ejercicio 1

```
import torch

print(torch.cuda.is_available())
```

True

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(device)
```

cuda

```
import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
from torchvision import datasets, transforms

# Definición del modelo
class LeNet5(nn.Module):
    def __init__(self):
        super(LeNet5, self).__init__()
        self.conv1 = nn.Conv2d(1, 6, 5, padding=2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16*5*5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = torch.tanh(F.avg_pool2d(self.conv1(x), 2))
        x = torch.tanh(F.avg_pool2d(self.conv2(x), 2))
        x = x.view(-1, 16*5*5)
        x = torch.sigmoid(self.fc1(x))
        x = torch.sigmoid(self.fc2(x))
        x = self.fc3(x)
        return F.log_softmax(x, dim=1)

# Hiperparámetros
BATCH_SIZE = 64
EPOCHS = 15
LR = 0.01
```

```

# Transformaciones y carga del dataset
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
])

train_dataset = datasets.MNIST('./data', train=True, download=True, transform=transform)
test_dataset = datasets.MNIST('./data', train=False, transform=transform)

train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=True)
test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=BATCH_SIZE, shuffle=False)

# Instancia del modelo, función de pérdida y optimizador
model = LeNet5().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=LR)

# Entrenamiento
for epoch in range(EPOCHS):
    model.train()
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        output = model(data)
        loss = criterion(output, target)
        loss.backward()
        optimizer.step()
        if batch_idx % 100 == 0:
            print(f"Train Epoch: {epoch} [{batch_idx*len(data)}/{len(train_loader.dataset)} "
                  f"({100. * batch_idx / len(train_loader):.0f}%)]\tLoss: {loss.item():.6f}")

# Evaluación
model.eval()
test_loss = 0
correct = 0
with torch.no_grad():
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        output = model(data)
        test_loss += criterion(output, target).item()
        pred = output.argmax(dim=1, keepdim=True)
        correct += pred.eq(target.view_as(pred)).sum().item()

print(f"({100. * correct / len(test_loader.dataset):.0f}%)")

```

```

Train Epoch: 0 [0/60000 (0%)]    Loss: 2.290354
Train Epoch: 0 [6400/60000 (11%)]    Loss: 2.315570
Train Epoch: 0 [12800/60000 (21%)]    Loss: 2.299078

```

Train Epoch: 0 [19200/60000 (32%)] Loss: 2.298903  
Train Epoch: 0 [25600/60000 (43%)] Loss: 2.300027  
Train Epoch: 0 [32000/60000 (53%)] Loss: 2.312724  
Train Epoch: 0 [38400/60000 (64%)] Loss: 2.307347  
Train Epoch: 0 [44800/60000 (75%)] Loss: 2.314235  
Train Epoch: 0 [51200/60000 (85%)] Loss: 2.299103  
Train Epoch: 0 [57600/60000 (96%)] Loss: 2.297892  
Train Epoch: 1 [0/60000 (0%)] Loss: 2.300854  
Train Epoch: 1 [6400/60000 (11%)] Loss: 2.304880  
Train Epoch: 1 [12800/60000 (21%)] Loss: 2.295555  
Train Epoch: 1 [19200/60000 (32%)] Loss: 2.306973  
Train Epoch: 1 [25600/60000 (43%)] Loss: 2.288721  
Train Epoch: 1 [32000/60000 (53%)] Loss: 2.300674  
Train Epoch: 1 [38400/60000 (64%)] Loss: 2.300735  
Train Epoch: 1 [44800/60000 (75%)] Loss: 2.296987  
Train Epoch: 1 [51200/60000 (85%)] Loss: 2.297271  
Train Epoch: 1 [57600/60000 (96%)] Loss: 2.303181  
Train Epoch: 2 [0/60000 (0%)] Loss: 2.297223  
Train Epoch: 2 [6400/60000 (11%)] Loss: 2.307899  
Train Epoch: 2 [12800/60000 (21%)] Loss: 2.305202  
Train Epoch: 2 [19200/60000 (32%)] Loss: 2.299099  
Train Epoch: 2 [25600/60000 (43%)] Loss: 2.278854  
Train Epoch: 2 [32000/60000 (53%)] Loss: 2.305628  
Train Epoch: 2 [38400/60000 (64%)] Loss: 2.292988  
Train Epoch: 2 [44800/60000 (75%)] Loss: 2.284897  
Train Epoch: 2 [51200/60000 (85%)] Loss: 2.301391  
Train Epoch: 2 [57600/60000 (96%)] Loss: 2.292528  
Train Epoch: 3 [0/60000 (0%)] Loss: 2.300167  
Train Epoch: 3 [6400/60000 (11%)] Loss: 2.318235  
Train Epoch: 3 [12800/60000 (21%)] Loss: 2.282169  
Train Epoch: 3 [19200/60000 (32%)] Loss: 2.308214  
Train Epoch: 3 [25600/60000 (43%)] Loss: 2.293639  
Train Epoch: 3 [32000/60000 (53%)] Loss: 2.287900  
Train Epoch: 3 [38400/60000 (64%)] Loss: 2.289978  
Train Epoch: 3 [44800/60000 (75%)] Loss: 2.289389  
Train Epoch: 3 [51200/60000 (85%)] Loss: 2.282530  
Train Epoch: 3 [57600/60000 (96%)] Loss: 2.280149  
Train Epoch: 4 [0/60000 (0%)] Loss: 2.293365  
Train Epoch: 4 [6400/60000 (11%)] Loss: 2.291999  
Train Epoch: 4 [12800/60000 (21%)] Loss: 2.289127  
Train Epoch: 4 [19200/60000 (32%)] Loss: 2.285772  
Train Epoch: 4 [25600/60000 (43%)] Loss: 2.282268  
Train Epoch: 4 [32000/60000 (53%)] Loss: 2.272134  
Train Epoch: 4 [38400/60000 (64%)] Loss: 2.267790  
Train Epoch: 4 [44800/60000 (75%)] Loss: 2.270984  
Train Epoch: 4 [51200/60000 (85%)] Loss: 2.259552  
Train Epoch: 4 [57600/60000 (96%)] Loss: 2.279252  
Train Epoch: 5 [0/60000 (0%)] Loss: 2.265924  
Train Epoch: 5 [6400/60000 (11%)] Loss: 2.263482  
Train Epoch: 5 [12800/60000 (21%)] Loss: 2.271164

Train Epoch: 5 [19200/60000 (32%)] Loss: 2.255219  
Train Epoch: 5 [25600/60000 (43%)] Loss: 2.227378  
Train Epoch: 5 [32000/60000 (53%)] Loss: 2.241599  
Train Epoch: 5 [38400/60000 (64%)] Loss: 2.201749  
Train Epoch: 5 [44800/60000 (75%)] Loss: 2.225660  
Train Epoch: 5 [51200/60000 (85%)] Loss: 2.214654  
Train Epoch: 5 [57600/60000 (96%)] Loss: 2.207725  
Train Epoch: 6 [0/60000 (0%)] Loss: 2.238866  
Train Epoch: 6 [6400/60000 (11%)] Loss: 2.173890  
Train Epoch: 6 [12800/60000 (21%)] Loss: 2.147619  
Train Epoch: 6 [19200/60000 (32%)] Loss: 2.129619  
Train Epoch: 6 [25600/60000 (43%)] Loss: 2.081799  
Train Epoch: 6 [32000/60000 (53%)] Loss: 2.104464  
Train Epoch: 6 [38400/60000 (64%)] Loss: 2.122292  
Train Epoch: 6 [44800/60000 (75%)] Loss: 2.079769  
Train Epoch: 6 [51200/60000 (85%)] Loss: 2.041947  
Train Epoch: 6 [57600/60000 (96%)] Loss: 1.935064  
Train Epoch: 7 [0/60000 (0%)] Loss: 1.980747  
Train Epoch: 7 [6400/60000 (11%)] Loss: 1.821163  
Train Epoch: 7 [12800/60000 (21%)] Loss: 1.860492  
Train Epoch: 7 [19200/60000 (32%)] Loss: 1.967731  
Train Epoch: 7 [25600/60000 (43%)] Loss: 1.738807  
Train Epoch: 7 [32000/60000 (53%)] Loss: 1.842051  
Train Epoch: 7 [38400/60000 (64%)] Loss: 1.723595  
Train Epoch: 7 [44800/60000 (75%)] Loss: 1.634346  
Train Epoch: 7 [51200/60000 (85%)] Loss: 1.704315  
Train Epoch: 7 [57600/60000 (96%)] Loss: 1.696681  
Train Epoch: 8 [0/60000 (0%)] Loss: 1.585786  
Train Epoch: 8 [6400/60000 (11%)] Loss: 1.460944  
Train Epoch: 8 [12800/60000 (21%)] Loss: 1.504637  
Train Epoch: 8 [19200/60000 (32%)] Loss: 1.538492  
Train Epoch: 8 [25600/60000 (43%)] Loss: 1.371103  
Train Epoch: 8 [32000/60000 (53%)] Loss: 1.488982  
Train Epoch: 8 [38400/60000 (64%)] Loss: 1.280876  
Train Epoch: 8 [44800/60000 (75%)] Loss: 1.339621  
Train Epoch: 8 [51200/60000 (85%)] Loss: 1.340497  
Train Epoch: 8 [57600/60000 (96%)] Loss: 1.295983  
Train Epoch: 9 [0/60000 (0%)] Loss: 1.223771  
Train Epoch: 9 [6400/60000 (11%)] Loss: 1.143815  
Train Epoch: 9 [12800/60000 (21%)] Loss: 1.285403  
Train Epoch: 9 [19200/60000 (32%)] Loss: 1.114302  
Train Epoch: 9 [25600/60000 (43%)] Loss: 1.140909  
Train Epoch: 9 [32000/60000 (53%)] Loss: 1.214866  
Train Epoch: 9 [38400/60000 (64%)] Loss: 1.105295  
Train Epoch: 9 [44800/60000 (75%)] Loss: 1.114371  
Train Epoch: 9 [51200/60000 (85%)] Loss: 0.957275  
Train Epoch: 9 [57600/60000 (96%)] Loss: 1.113677  
Train Epoch: 10 [0/60000 (0%)] Loss: 1.016166  
Train Epoch: 10 [6400/60000 (11%)] Loss: 0.836244  
Train Epoch: 10 [12800/60000 (21%)] Loss: 0.917692

Train Epoch: 10 [19200/60000 (32%)] Loss: 0.962960  
Train Epoch: 10 [25600/60000 (43%)] Loss: 0.914223  
Train Epoch: 10 [32000/60000 (53%)] Loss: 0.935577  
Train Epoch: 10 [38400/60000 (64%)] Loss: 0.927864  
Train Epoch: 10 [44800/60000 (75%)] Loss: 0.908553  
Train Epoch: 10 [51200/60000 (85%)] Loss: 0.985946  
Train Epoch: 10 [57600/60000 (96%)] Loss: 0.797681  
Train Epoch: 11 [0/60000 (0%)] Loss: 0.778887  
Train Epoch: 11 [6400/60000 (11%)] Loss: 0.921462  
Train Epoch: 11 [12800/60000 (21%)] Loss: 0.739085  
Train Epoch: 11 [19200/60000 (32%)] Loss: 0.915735  
Train Epoch: 11 [25600/60000 (43%)] Loss: 0.749777  
Train Epoch: 11 [32000/60000 (53%)] Loss: 0.927942  
Train Epoch: 11 [38400/60000 (64%)] Loss: 0.868798  
Train Epoch: 11 [44800/60000 (75%)] Loss: 0.705947  
Train Epoch: 11 [51200/60000 (85%)] Loss: 0.867870  
Train Epoch: 11 [57600/60000 (96%)] Loss: 0.877646  
Train Epoch: 12 [0/60000 (0%)] Loss: 0.821957  
Train Epoch: 12 [6400/60000 (11%)] Loss: 0.677227  
Train Epoch: 12 [12800/60000 (21%)] Loss: 0.888611  
Train Epoch: 12 [19200/60000 (32%)] Loss: 0.587170  
Train Epoch: 12 [25600/60000 (43%)] Loss: 0.953355  
Train Epoch: 12 [32000/60000 (53%)] Loss: 0.589124  
Train Epoch: 12 [38400/60000 (64%)] Loss: 0.667878  
Train Epoch: 12 [44800/60000 (75%)] Loss: 0.710181  
Train Epoch: 12 [51200/60000 (85%)] Loss: 0.603574  
Train Epoch: 12 [57600/60000 (96%)] Loss: 0.790573  
Train Epoch: 13 [0/60000 (0%)] Loss: 0.381692  
Train Epoch: 13 [6400/60000 (11%)] Loss: 0.691773  
Train Epoch: 13 [12800/60000 (21%)] Loss: 0.644380  
Train Epoch: 13 [19200/60000 (32%)] Loss: 0.597041  
Train Epoch: 13 [25600/60000 (43%)] Loss: 0.514326  
Train Epoch: 13 [32000/60000 (53%)] Loss: 0.674853  
Train Epoch: 13 [38400/60000 (64%)] Loss: 0.462263  
Train Epoch: 13 [44800/60000 (75%)] Loss: 0.619597  
Train Epoch: 13 [51200/60000 (85%)] Loss: 0.621748  
Train Epoch: 13 [57600/60000 (96%)] Loss: 0.470695  
Train Epoch: 14 [0/60000 (0%)] Loss: 0.660009  
Train Epoch: 14 [6400/60000 (11%)] Loss: 0.599284  
Train Epoch: 14 [12800/60000 (21%)] Loss: 0.528506  
Train Epoch: 14 [19200/60000 (32%)] Loss: 0.608190  
Train Epoch: 14 [25600/60000 (43%)] Loss: 0.532200  
Train Epoch: 14 [32000/60000 (53%)] Loss: 0.468184  
Train Epoch: 14 [38400/60000 (64%)] Loss: 0.483214  
Train Epoch: 14 [44800/60000 (75%)] Loss: 0.643663  
Train Epoch: 14 [51200/60000 (85%)] Loss: 0.454628  
Train Epoch: 14 [57600/60000 (96%)] Loss: 0.467120  
(87%)

## Ejercicio 2

```
import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transforms
import torch.nn.functional as F

class AlexNet(nn.Module):
    def __init__(self, num_classes=1000):
        super(AlexNet, self).__init__()

        # Capa convolucional 1
        self.conv1 = nn.Conv2d(3, 96, kernel_size=11, stride=4, padding=2)
        self.relu1 = nn.ReLU(inplace=True)
        self.lrn1 = nn.LocalResponseNorm(5, alpha=1e-4, beta=0.75, k=2)
        self.pool1 = nn.MaxPool2d(kernel_size=3, stride=2)

        # Capa convolucional 2
        self.conv2 = nn.Conv2d(96, 256, kernel_size=5, stride=1, padding=2)
        self.relu2 = nn.ReLU(inplace=True)
        self.lrn2 = nn.LocalResponseNorm(5, alpha=1e-4, beta=0.75, k=2)
        self.pool2 = nn.MaxPool2d(kernel_size=3, stride=2)

        # Capa convolucional 3
        self.conv3 = nn.Conv2d(256, 384, kernel_size=3, stride=1, padding=1)
        self.relu3 = nn.ReLU(inplace=True)

        # Capa convolucional 4
        self.conv4 = nn.Conv2d(384, 384, kernel_size=3, stride=1, padding=1)
        self.relu4 = nn.ReLU(inplace=True)

        # Capa convolucional 5
        self.conv5 = nn.Conv2d(384, 256, kernel_size=3, stride=1, padding=1)
        self.relu5 = nn.ReLU(inplace=True)
        self.pool5 = nn.MaxPool2d(kernel_size=3, stride=2)

        # Capas completamente conectadas
        self.fc6 = nn.Linear(256 * 6 * 6, 4096)
        self.relu6 = nn.ReLU(inplace=True)
        self.dropout6 = nn.Dropout(0.5)

        self.fc7 = nn.Linear(4096, 4096)
        self.relu7 = nn.ReLU(inplace=True)
        self.dropout7 = nn.Dropout(0.5)

        self.fc8 = nn.Linear(4096, num_classes)

    def forward(self, x):
        x = self.pool1(self.lrn1(self.relu1(self.conv1(x))))
        x = self.pool2(self.lrn2(self.relu2(self.conv2(x))))
```

```

        x = self.relu3(self.conv3(x))
        x = self.relu4(self.conv4(x))
        x = self.pool5(self.relu5(self.conv5(x)))

        x = x.view(x.size(0), 256 * 6 * 6)

        x = self.dropout6(self.relu6(self.fc6(x)))
        x = self.dropout7(self.relu7(self.fc7(x)))

        x = self.fc8(x)

        return F.softmax(x, dim=1)

# Función para entrenar el modelo
def train_model(model, train_loader, criterion, optimizer, num_epochs=30):
    for epoch in range(num_epochs):
        model.train()
        running_loss = 0.0
        for images, labels in train_loader:
            images, labels = images.to(device), labels.to(device)
            optimizer.zero_grad()
            outputs = model(images)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()
            running_loss += loss.item()
        print(f"Epoch {epoch+1}/{num_epochs}, Loss: {running_loss/len(train_loader)}")

# Función para evaluar el modelo
def evaluate_model(model, test_loader):
    model.eval()
    correct = 0
    total = 0
    with torch.no_grad():
        for images, labels in test_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            _, predicted = outputs.max(1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()
    accuracy = 100 * correct / total
    print(f"Accuracy on test set: {accuracy}%")

transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

train_dataset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=
```

```

test_dataset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, transform=

train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=64, shuffle=True)
test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=64, shuffle=False)

model = AlexNet(num_classes=10).to(device)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=0.01, momentum=0.9)

train_model(model, train_loader, criterion, optimizer, num_epochs=30)

evaluate_model(model, test_loader)

```

Files already downloaded and verified

Files already downloaded and verified

Epoch 1/30, Loss: 2.302589490895381  
 Epoch 2/30, Loss: 2.3025704125309234  
 Epoch 3/30, Loss: 2.3025254858729176  
 Epoch 4/30, Loss: 2.302226912944823  
 Epoch 5/30, Loss: 2.25119497983352  
 Epoch 6/30, Loss: 2.213179388619445  
 Epoch 7/30, Loss: 2.1572559105465783  
 Epoch 8/30, Loss: 2.111590341533846  
 Epoch 9/30, Loss: 2.071355868635885  
 Epoch 10/30, Loss: 2.041391796163281  
 Epoch 11/30, Loss: 1.9850449238896675  
 Epoch 12/30, Loss: 1.9442216006996076  
 Epoch 13/30, Loss: 1.9191369117068513  
 Epoch 14/30, Loss: 1.886217120815726  
 Epoch 15/30, Loss: 1.8580890584479817  
 Epoch 16/30, Loss: 1.833253474948961  
 Epoch 17/30, Loss: 1.8073494653872517  
 Epoch 18/30, Loss: 1.7937313595696178  
 Epoch 19/30, Loss: 1.7717730893808252  
 Epoch 20/30, Loss: 1.7616859060114303  
 Epoch 21/30, Loss: 1.758428267048448  
 Epoch 22/30, Loss: 1.7453365409770585  
 Epoch 23/30, Loss: 1.732283755031693  
 Epoch 24/30, Loss: 1.710803150216027  
 Epoch 25/30, Loss: 1.705536770546223  
 Epoch 26/30, Loss: 1.6985302739740942  
 Epoch 27/30, Loss: 1.695783346967624  
 Epoch 28/30, Loss: 1.6854584367988665  
 Epoch 29/30, Loss: 1.678006345658656  
 Epoch 30/30, Loss: 1.6738292587077832  
 Accuracy on test set: 75.48%

a. ¿Cuál es la diferencia principal entre ambas arquitecturas?



- LaNet-5: es una arquitectura más simple, contando con 2 capas convolucionales y 3 capas completamente conectadas. Se utiliza la función de activación de tanh y existe un menor número de parámetros y menor profundidad que logra AlexNet.
- AlexNet: es una arquitectura más compleja, que posee 5 capas convolucionales y 3 capas completamente conectadas. Utiliza la función de activación ReLU e implementa funciones como dropout para evitar el sobreajuste. Permite mayor profundidad y tiene un mayor número de parámetros que LaNet-5.

b. Podría usarse LeNet-5 para un problema como el que resolvió usando AlexNet? ¿Y viceversa?

- Técnicamente LeNet-5 si se podría usar para problemas que resolvió Alexnet. Sin embargo LeNet-5 tiene un menor capacidad y profundidad haciendo que no funcione tan bien como Alexnet. Por otro lado, Alexnet si se podría usar para problemas que resolvió LeNet-5, pero al tener una mayor capacidad y profundidad. Esta mayor capacidad puede hacer que se use un modelo más complejo para problemas más sencillos, lo cual puede llevar a un uso innecesario de recursos.

c. Indique claramente qué le pareció más interesante de cada arquitectura

- Lo que más nos llamó la atención de LeNet-5 es la simplicidad del modelo. A pesar de ser considerablemente simple, logró obtener resultados bastante buenos. Por otro lado, lo que más nos llamó la atención de Alexnet es la complejidad del modelo. A pesar de ser un modelo complejo, no logró resultados tan buenos como LeNet-5. Esto nos hace pensar que la complejidad de un modelo no necesariamente se traduce en mejores resultados.

Investigue e indique en qué casos son útiles las siguientes arquitecturas, agregue imágenes si esto le ayuda a una mejor comprensión

a. GoogleNet (Inception)

- GoogleNet, también conocida como Inception, es una arquitectura de CNN desarrollada por Google. Se destacó por su profundidad y eficiencia en la utilización de los recursos.
- Es útil en casos donde se requieren redes profundas pero se desea mantener un uso eficiente de los recursos computacionales. GoogleNet utiliza una estructura llamada "módulos Inception" que combina múltiples tamaños de filtros de convolución en paralelo, permitiendo la extracción de características a diferentes escalas.
- Útil para tareas de clasificación de imágenes, detección de objetos y segmentación semántica.

b. DenseNet (Densely Connected Convolutional Networks)

- DenseNet es una arquitectura de CNN que se caracteriza por su densa conectividad entre capas. Cada capa está conectada directamente con todas las capas subsiguientes.
- Es útil en casos donde se desea un mejor flujo de información y gradientes más fuertes a lo largo de la red, lo que facilita el entrenamiento de redes profundas.
- Útil para tareas de clasificación de imágenes, detección de objetos y segmentación semántica.

c. MobileNet

- MobileNet es una arquitectura de CNN diseñada para aplicaciones en dispositivos móviles y embebidos con recursos computacionales limitados.
- Es útil en casos donde se necesita una red ligera y rápida, como en aplicaciones de visión por computadora en dispositivos móviles.
- Útil para tareas de clasificación de imágenes, detección de objetos en tiempo real y otras aplicaciones de visión en dispositivos móviles.

#### d. EfficientNet

- EfficientNet es una familia de arquitecturas de CNN que buscan optimizar el equilibrio entre el rendimiento y la eficiencia computacional mediante el uso de escalado compuesto.
- Es útil en casos donde se desean modelos con un buen rendimiento pero que sean escalables en términos de tamaño y requisitos computacionales.
- Útil para una variedad de tareas de visión por computadora, desde clasificación de imágenes hasta detección de objetos y segmentación semántica.

¿Cómo la arquitectura de transformers puede ser usada para image recognition?

La arquitectura de Transformers se puede usar en el reconocimiento de imágenes al tratar las imágenes como secuencias de parches y aplicar mecanismos de atención y transformación para capturar información espacial y contextual. Esto se puede hacer al obtener las características, el uso de atención multi-cabeza y agregar información. Los modelos de visión Transformer se entrenan en conjuntos de datos etiquetados, se ajustan finamente en tareas específicas y permiten la transferencia de aprendizaje.